

EDA and data visualization

Monica Alexander

01/02/23

Table of contents

1	Overview	1
1.1	What to hand in via GitHub	2
1.2	A note on packages	2
2	TTC subway delays	3
3	EDA and data viz	4
3.1	Data checks	5
3.1.1	Sanity Checks	5
3.1.2	Missing values	6
3.1.3	Duplicates?	7
3.2	Visualizing distributions	7
3.3	Visualizing time series	13
3.4	Visualizing relationships	15
3.5	PCA (additional)	16
4	Lab Exercises	19

1 Overview

This week we will be going through some exploratory data analysis (EDA) and data visualization steps in R. The aim is to get you used to working with real data (that has issues) to understand the main characteristics and potential issues.

We will be using the [opendatatoronto](#) R package, which interfaces with the City of Toronto Open Data Portal.

A good resource is part 1 (especially chapters 3 and 7) of ‘R for Data Science’ by Hadley Wickham, available for free here: <https://r4ds.had.co.nz/>.

1.1 What to hand in via GitHub

There are exercises at the end of this lab. Please make a new .Rmd file with your answers, call it something sensible (e.g. `week_2_lab.Rmd`), commit to your git repo from last week, and push to GitHub. Due on Monday by 9am.

1.2 A note on packages

You may need to install various packages used (using the `install.packages` function). Load in all the packages we need:

```
library(opendatatoronto)
```

Warning: package 'opendatatoronto' was built under R version 4.2.2

```
library(tidyverse)
```

Warning: package 'tidyverse' was built under R version 4.2.2

```
library(stringr)
library(skimr) # EDA
```

Warning: package 'skimr' was built under R version 4.2.2

```
library(visdat) # EDA
```

Warning: package 'visdat' was built under R version 4.2.2

```
library(janitor)
```

Warning: package 'janitor' was built under R version 4.2.2

```
library(lubridate)
```

Warning: package 'lubridate' was built under R version 4.2.2

Warning: package 'timechange' was built under R version 4.2.2

```
library(ggrepel)
```

Warning: package 'ggrepel' was built under R version 4.2.2

2 TTC subway delays

This package provides an interface to all data available on the [Open Data Portal](#) provided by the City of Toronto.

Use the `list_packages` function to look what's available

```
all_data <- list_packages(limit = 500)
head(all_data)
```

```
# A tibble: 6 x 11
  title      id    topics civic~1 publi~2 excerpt datas~3 num_r~4 formats refre~5
  <chr>    <chr> <chr>  <chr>  <chr>  <chr>  <chr>    <int> <chr>  <chr>
1 Developm~ 0aa7~ <NA>   <NA>   City P~ "This ~ Table      4 CSV,XM~ Monthly
2 Polls co~ 7bce~ City ~ <NA>   City C~ "Polls~ Table      5 XML,JS~ Daily
3 COVID-19~ d3f2~ Health <NA>   Toront~ "This ~ Map      13 GEOJSO~ Daily
4 Toronto'~ c6d6~ <NA>   <NA>   City M~ "This ~ Table      4 XML,JS~ Daily
5 Committe~ 260e~ City ~ Afford~ City P~ "This ~ Table     96 CSV,XM~ Weekly
6 Apartmen~ 4ef8~ Locat~ Afford~ Munici~ "This ~ Table      4 XML,JS~ Daily
# ... with 1 more variable: last_refreshed <date>, and abbreviated variable
#   names 1: civic_issues, 2: publisher, 3: dataset_category, 4: num_resources,
#   5: refresh_rate
```

Let's download the data on TTC subway delays in 2022.

```
res <- list_package_resources("996cfe8d-fb35-40ce-b569-698d51fc683b") # obtained code from
res <- res |> mutate(year = str_extract(name, "202.?"))
delay_2022_ids <- res |> filter(year==2022) |> select(id) |> pull()

delay_2022 <- get_resource(delay_2022_ids)

# make the column names nicer to work with
delay_2022 <- clean_names(delay_2022)
```

Let's also download the delay code and readme, as reference.

```
# note: I obtained these codes from the 'id' column in the `res` object above
delay_codes <- get_resource("3900e649-f31e-4b79-9f20-4731bbfd94f7")
```

New names:

```
* `` -> `...1`
* `CODE DESCRIPTION` -> `CODE DESCRIPTION...3`
* `` -> `...4`
* `` -> `...5`
* `CODE DESCRIPTION` -> `CODE DESCRIPTION...7`
```

```
delay_data_codebook <- get_resource("ca43ac3d-3940-4315-889b-a9375e7b8aa4")
```

This dataset has a bunch of interesting variables. You can refer to the readme for descriptions. Our outcome of interest is `min_delay`, which give the delay in mins.

```
head(delay_2022)
```

```
# A tibble: 6 x 10
  date          time day      station  code min_d~1 min_gap bound line
<dtm>         <chr> <chr>   <chr>   <chr>   <dbl>   <dbl> <chr> <chr>
1 2022-01-01 00:00:00 15:59 Saturday LAWRENCE~ SRDP      0      0 N     SRT
2 2022-01-01 00:00:00 02:23 Saturday SPADINA ~ MUIS      0      0 <NA> BD
3 2022-01-01 00:00:00 22:00 Saturday KENNEDY ~ MRO      0      0 <NA> SRT
4 2022-01-01 00:00:00 02:28 Saturday VAUGHAN ~ MUIS      0      0 <NA> YU
5 2022-01-01 00:00:00 02:34 Saturday EGLINTON~ MUATC      0      0 S     YU
6 2022-01-01 00:00:00 05:40 Saturday QUEEN ST~ MUNCA      0      0 <NA> YU
# ... with 1 more variable: vehicle <dbl>, and abbreviated variable name
#   1: min_delay
```

3 EDA and data viz

The following section highlights some tools that might be useful for you when you are getting used to a new dataset. There's no one way of exploration, but it's important to always keep in mind:

- what should your variables look like (type, values, distribution, etc)
- what would be surprising (outliers etc)
- what is your end goal (here, it might be understanding factors associated with delays, e.g. stations, time of year, time of day, etc)

In any data analysis project, if it turns out you have data issues, surprising values, missing data etc, it's important you **document** anything you found and the subsequent steps or **assumptions** you made before moving onto your data analysis / modeling.

3.1 Data checks

3.1.1 Sanity Checks

We need to check variables should be what they say they are. If they aren't, the natural next question is to what to do with issues (recode? remove?)

E.g. check days of week

```
unique(delay_2022$day)
```

```
[1] "Saturday" "Sunday"    "Monday"    "Tuesday"   "Wednesday" "Thursday"
[7] "Friday"
```

Check lines: oh no. some issues here. Some have obvious recodes, others, not so much.

```
unique(delay_2022$line)
```

```
[1] "SRT"          "BD"           "YU"           "YU/BD"
[5] "SHP"          NA             "BD/YU"        "YU / BD"
[9] "YU/ BD"       "B/D"          "Y/BD"         "YU/BD LINES"
[13] "YUS"          "YU & BD"      "YUS AND BD"   "YUS/BD"
[17] "69 WARDEN SOUTH" "YU/BD LINE"  "LINE 2 SHUTTLE" "57 MIDLAND"
[21] "96 WILSON"     "506 CARLTON"
```

The `skimr` package might also be useful here

```
skim(delay_2022)
```

Table 1: Data summary

Name	delay_2022
Number of rows	19895
Number of columns	10

Table 1: Data summary

Column type frequency:	
character	6
numeric	3
POSIXct	1
Group variables	
None	

Variable type: character

skim_variable	n_missing	complete_rate	min	max	empty	n_unique	whitespace
time	0	1.00	5	5	0	1406	0
day	0	1.00	6	9	0	7	0
station	0	1.00	5	22	0	296	0
code	0	1.00	3	5	0	179	0
bound	5546	0.72	1	1	0	5	0
line	39	1.00	2	15	0	21	0

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
min_delay	0	1	3.67	12.00	0	0	0	4	458	
min_gap	0	1	5.33	12.66	0	0	0	8	463	
vehicle	0	1	3571.59	2646.62	0	0	5192	5701	8871	

Variable type: POSIXct

skim_variable	n_missing	complete_rate	min	max	median	n_unique
date	0	1	2022-01-01	2022-12-31	2022-06-29	365

3.1.2 Missing values

Calculate number of NAs by column

```
delay_2022 |>
  summarize(across(everything(), ~ sum(is.na(.x))))
```

```
# A tibble: 1 x 10
  date   time   day station  code min_delay min_gap bound  line vehicle
<int> <int> <int> <int> <int>    <int>    <int> <int> <int>    <int>
1     0     0     0       0     0        0        0 5546    39      0
```

The `visdat` package is useful here, particularly to see how missing values are distributed. (commented out because couldn't get pdf to render in quarto)

```
#vis_dat(delay_2022)
#vis_miss(delay_2022)
```

3.1.3 Duplicates?

The `get_dupes` function from the `janitor` package is useful for this.

```
get_dupes(delay_2022)
```

No variable names specified - using all columns.

```
# A tibble: 28 x 11
  date           time day      station code min_d~1 min_gap bound line
<dtm>          <chr> <chr>    <chr>  <chr>    <dbl>    <dbl> <chr> <chr>
1 2022-01-12 00:00:00 13:27 Wednesday FINCH ~ TUNOA      3        6 S    YU
2 2022-01-12 00:00:00 13:27 Wednesday FINCH ~ TUNOA      3        6 S    YU
3 2022-01-12 00:00:00 17:49 Wednesday FINCH ~ TUNOA      3        6 S    YU
4 2022-01-12 00:00:00 17:49 Wednesday FINCH ~ TUNOA      3        6 S    YU
5 2022-01-17 00:00:00 02:00 Monday    SCARBO~ TRST      0         0 <NA> SRT
6 2022-01-17 00:00:00 02:00 Monday    SCARBO~ TRST      0         0 <NA> SRT
7 2022-01-20 00:00:00 02:30 Thursday YONGE ~ TUST      0         0 <NA> YU
8 2022-01-20 00:00:00 02:30 Thursday YONGE ~ TUST      0         0 <NA> YU
9 2022-01-20 00:00:00 08:51 Thursday WILSON~ TUNOA      3         6 S    YU
10 2022-01-20 00:00:00 08:51 Thursday WILSON~ TUNOA      3         6 S    YU
# ... with 18 more rows, 2 more variables: vehicle <dbl>, dupe_count <int>, and
# abbreviated variable name 1: min_delay
```

3.2 Visualizing distributions

Histograms, barplots, and density plots are your friends here.


Let's look at the outcome of interest: `min_delay`. First of all just a histogram of all the data:

```
## Removing the observations that have non-standardized lines

delay_2022 <- delay_2022 |> filter(line %in% c("BD", "YU", "SHP", "SRT"))

ggplot(data = delay_2022) +
  geom_histogram(aes(x = min_delay))
```

`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.




To improve readability, could plot on logged scale:

```
ggplot(data = delay_2022) +
  geom_histogram(aes(x = min_delay)) +
  scale_x_log10()
```

Warning: Transformation introduced infinite values in continuous x-axis

`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.

Warning: Removed 10500 rows containing non-finite values (stat_bin).



Our initial EDA hinted at an outlying delay time, let's take a look at the largest delays below. Join the `delay_codes` dataset to see what the delay is. (Have to do some mangling as SRT has different codes).

```
delay_2022 <- delay_2022 |>
  left_join(delay_codes |> rename(code = `SUB RMENU CODE`, code_desc = `CODE DESCRIPTION..`))
```

Joining, by = "code"

```
delay_2022 <- delay_2022 |>
  mutate(code_srt = ifelse(line=="SRT", code, "NA")) |>
  left_join(delay_codes |> rename(code_srt = `SRT RMENU CODE`, code_desc_srt = `CODE DESCRIPTION..`)) |>
  mutate(code = ifelse(code_srt=="NA", code, code_srt),
         code_desc = ifelse(is.na(code_desc_srt), code_desc, code_desc_srt)) |>
  select(-code_srt, -code_desc_srt)
```

Joining, by = "code_srt"

The largest delay is due to "Signals Other".

```

delay_2022 |>
  left_join(delay_codes |> rename(code = `SUB RMENU CODE`, code_desc = `CODE DESCRIPTION..
  arrange(-min_delay) |>
  select(date, time, station, line, min_delay, code, code_desc)

```

Joining, by = c("code", "code_desc")

A tibble: 19,473 x 7

	date	time	station	line	min_de~1	code	code_~2
	<dtm>	<chr>	<chr>	<chr>	<dbl>	<chr>	<chr>
1	2022-12-08 00:00:00	17:52	MIDLAND STATION	SRT	458	MRPLB	Fire/S~
2	2022-08-22 00:00:00	12:20	SRT LINE	SRT	451	PRSO	Signal~
3	2022-04-28 00:00:00	06:02	JANE STATION	BD	388	PUTR	Rail R~
4	2022-07-26 00:00:00	07:06	YONGE BD STATION	BD	382	MUPLB	Fire/S~
5	2022-08-15 00:00:00	12:57	DUFFERIN STATION	BD	327	MUPR1	Priori~
6	2022-01-26 00:00:00	20:15	KENNEDY SRT STATION	SRT	315	MRWEA	Weathe~
7	2022-08-02 00:00:00	21:23	HIGHWAY 407 STATION	YU	312	MUPR1	Priori~
8	2022-01-17 00:00:00	21:30	SHEPPARD WEST TO UNION	YU	291	MUFM	Force ~
9	2022-01-25 00:00:00	21:03	SCARBOROUGH CTR STATIO	SRT	285	PRSL	Loop R~
10	2022-06-17 00:00:00	12:25	KIPLING STATION	BD	241	SUUT	Unauth~

... with 19,463 more rows, and abbreviated variable names 1: min_delay,
2: code_desc

3.2.0.1 Grouping and small multiples

A quick and powerful visualization technique is to group the data by a variable of interest, e.g. line


```

ggplot(data = delay_2022) +
  geom_histogram(aes(x = min_delay, y = ..density.., fill = line), position = 'dodge', bin
  scale_x_log10()

```

Warning: Transformation introduced infinite values in continuous x-axis

Warning: Removed 10500 rows containing non-finite values (stat_bin).




I switched to density above to look at the the distributions more comparably, but we should also be aware of differences in frequency, in particular, SHP and SRT have much smaller counts:

```
ggplot(data = delay_2022) +
  geom_histogram(aes(x = min_delay, fill = line), position = 'dodge', bins = 10) +
  scale_x_log10()
```

Warning: Transformation introduced infinite values in continuous x-axis

Warning: Removed 10500 rows containing non-finite values (stat_bin).

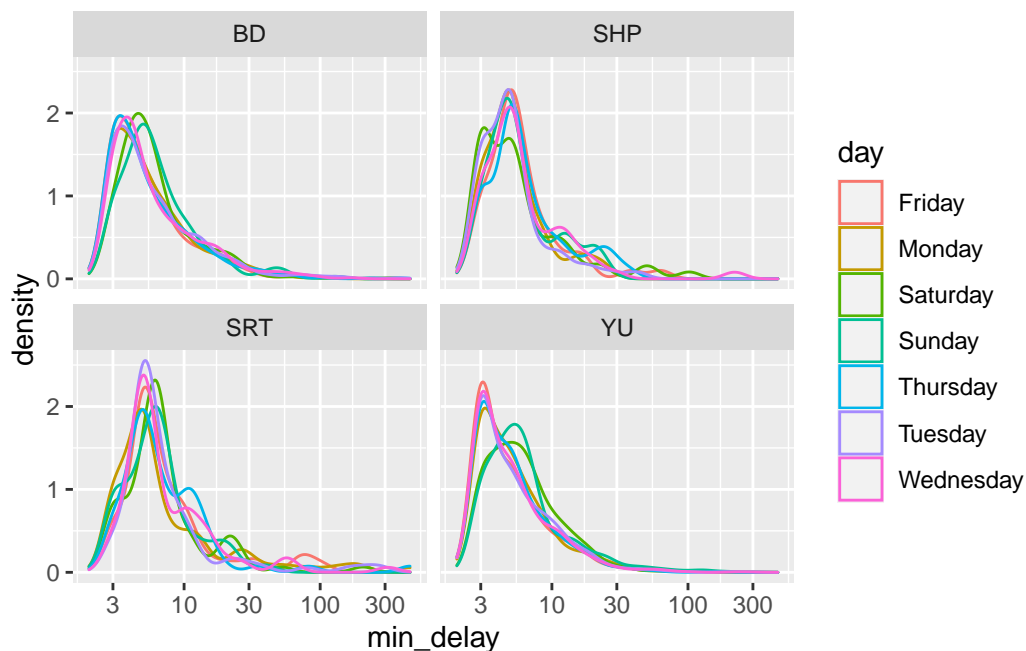


If you want to group by more than one variable, facets are good:

```
ggplot(data = delay_2022) +  
  geom_density(aes(x = min_delay, color = day), bw = .08) +  
  scale_x_log10() +  
  facet_wrap(~line)
```

Warning: Transformation introduced infinite values in continuous x-axis

Warning: Removed 10500 rows containing non-finite values (stat_density).



Side note: the station names are a mess. Try and clean up the station names a bit by taking just the first word (or, the first two if it starts with “ST”):


```
delay_2022 <- delay_2022 |>
  mutate(station_clean = ifelse(str_starts(station, "ST"), word(station, 1,2), word(station, 1)))
```

3.3 Visualizing time series

Daily plot is messy (you can check for yourself). Let’s look by week to see if there’s any seasonality. The `lubridate` package has lots of helpful functions that deal with date variables. First, mean delay (of those that were delayed more than 0 mins):

```
delay_2022 |>
  filter(min_delay>0) |>
  mutate(week = week(date)) |>
  group_by(week, line) |>
  summarise(mean_delay = mean(min_delay)) |>
  ggplot(aes(week, mean_delay, color = line)) +
  geom_point() +
  geom_smooth() +
  facet_grid(~line)
```


``summarise()`` has grouped output by 'week'. You can override using the ``.groups`` argument.
``geom_smooth()`` using `method = 'loess'` and formula `'y ~ x'`



What about proportion of delays that were greater than 10 mins?

```
delay_2022 |>
  mutate(week = week(date)) |>
  group_by(week, line) |>
  summarise(prop_delay = sum(min_delay>10)/n()) |>
  ggplot(aes(week, prop_delay, color = line)) +
  geom_point() +
  geom_smooth() +
  facet_grid(~line)
```

``summarise()`` has grouped output by 'week'. You can override using the ``.groups`` argument.
``geom_smooth()`` using `method = 'loess'` and formula `'y ~ x'`



3.4 Visualizing relationships

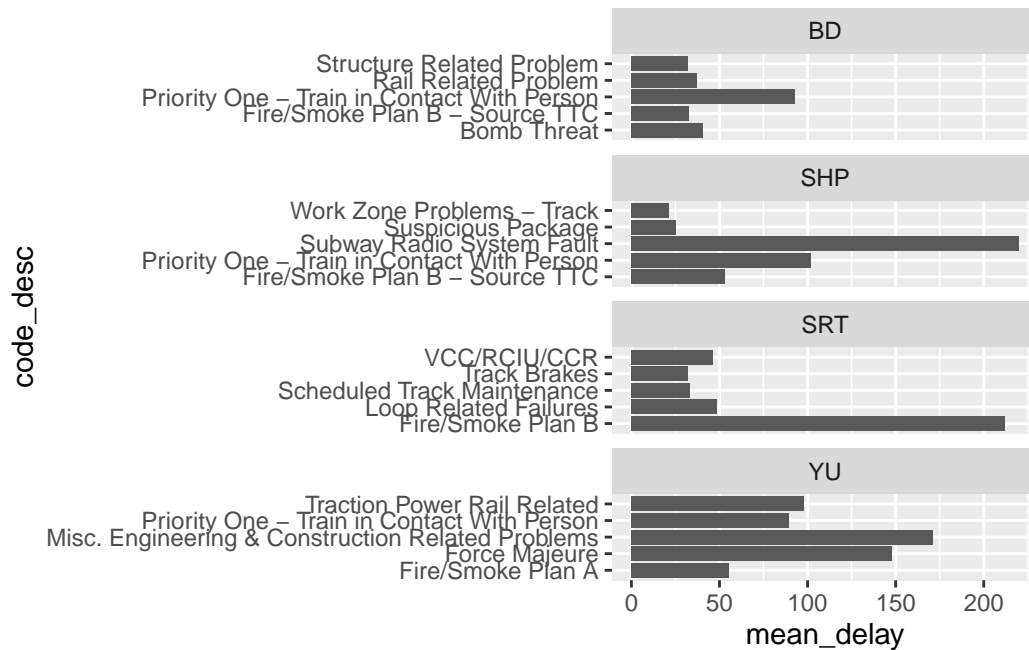
Note that **scatter plots** are a good precursor to modeling, to visualize relationships between continuous variables. Nothing obvious to plot here, but easy to do with `geom_point`.

Look at top five reasons for delay by station. Do they differ? Think about how this could be modeled.

```
delay_2022 |>
  group_by(line, code_desc) |>
  summarise(mean_delay = mean(min_delay)) |>
  arrange(-mean_delay) |>
  slice(1:5) |>
  ggplot(aes(x = code_desc,
             y = mean_delay)) +
  geom_col() +
  facet_wrap(vars(line),
            scales = "free_y",
            nrow = 4) +
  coord_flip()
```

``summarise()`` has grouped output by 'line'. You can override using the

` .groups` argument.



3.5 PCA (additional)

Principal components analysis is a really powerful exploratory tool, particularly when you have a lot of variables. It allows you to pick up potential clusters and/or outliers that can help to inform model building.

Let's do a quick (and imperfect) example looking at types of delays by station.

The delay categories are a bit of a mess, and there's hundreds of them. As a simple start, let's just take the first word:

```
delay_2022 <- delay_2022 |>
  mutate(code_red = case_when(
    str_starts(code_desc, "No") ~ word(code_desc, 1, 2),
    str_starts(code_desc, "Operator") ~ word(code_desc, 1,2),
    TRUE ~ word(code_desc,1))
  )
```

Let's also just restrict the analysis to causes that happen at least 50 times over 2022 To do the PCA, the dataframe also needs to be switched to wide format:


```

dwide <- delay_2022 |>
  group_by(line, station_clean) |>
  mutate(n_obs = n()) |>
  filter(n_obs>1) |>
  group_by(code_red) |>
  mutate(tot_delay = n()) |>
  arrange(tot_delay) |>
  filter(tot_delay>50) |>
  group_by(line, station_clean, code_red) |>
  summarise(n_delay = n()) |>
  pivot_wider(names_from = code_red, values_from = n_delay) |>
  mutate(
    across(everything(), ~ replace_na(.x, 0))
  )

```

`summarise()` has grouped output by 'line', 'station_clean'. You can override using the `.groups` argument.

Do the PCA:

```

delay_pca <- prcomp(dwide[,3:ncol(dwide)])

df_out <- as_tibble(delay_pca$x)
df_out <- bind_cols(dwide |> select(line, station_clean), df_out)
head(df_out)

```

```


# A tibble: 6 x 41
# Groups:   line, station_clean [6]
  line station~1 PC1 PC2 PC3 PC4 PC5 PC6 PC7 PC8 PC9
  <chr> <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 BD BATHURST -16.3 -24.3 -6.41 10.6 -3.15 5.36 -3.27 -8.68 -10.2
2 BD BAY 8.50 -13.2 -6.28 8.05 0.804 0.401 6.09 0.311 -1.66
3 BD BLOOR 36.3 28.7 34.2 14.7 9.70 7.70 -4.51 1.13 -2.31
4 BD BLOOR-DA~ 48.8 6.38 -0.483 1.44 -9.26 3.75 -0.671 0.388 0.189
5 BD BROADVIEW -22.6 -26.1 -6.18 11.8 4.18 -2.66 4.48 7.57 8.27
6 BD CASTLE 15.9 -8.44 -3.21 6.64 -3.65 0.366 0.624 -3.89 -2.99
# ... with 30 more variables: PC10 <dbl>, PC11 <dbl>, PC12 <dbl>, PC13 <dbl>,
# PC14 <dbl>, PC15 <dbl>, PC16 <dbl>, PC17 <dbl>, PC18 <dbl>, PC19 <dbl>,
# PC20 <dbl>, PC21 <dbl>, PC22 <dbl>, PC23 <dbl>, PC24 <dbl>, PC25 <dbl>,
# PC26 <dbl>, PC27 <dbl>, PC28 <dbl>, PC29 <dbl>, PC30 <dbl>, PC31 <dbl>,
# PC32 <dbl>, PC33 <dbl>, PC34 <dbl>, PC35 <dbl>, PC36 <dbl>, PC37 <dbl>,

```

```
# PC38 <dbl>, PC39 <dbl>, and abbreviated variable name 1: station_clean
```

Plot the first two PCs, and label some outlying stations:

```
ggplot(df_out,aes(x=PC1,y=PC2,color=line )) + geom_point() + geom_text_repel(data = df_out
```



Plot the factor loadings. Some evidence of public v operator?

```
df_out_r <- as_tibble(delay_pca$rotation)
df_out_r$feature <- colnames(dwide[,3:ncol(dwide)])

df_out_r
```

```
# A tibble: 39 x 40
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8	PC9
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	-0.127	-0.0381	-0.0174	0.0271	0.0387	-0.0425	0.122	-0.0238	0.159
2	-0.305	-0.127	-0.0743	0.0461	0.103	-0.183	0.190	-0.647	-0.493
3	-0.0530	-0.0113	0.0380	0.0382	0.0573	-0.0460	-0.0608	-0.116	0.250
4	-0.0135	-0.0171	-0.0117	-0.00271	0.0454	-0.0367	0.0137	0.0191	-0.0712
5	-0.0119	-0.00470	0.000218	0.00865	-0.0173	-0.0471	-0.0315	-0.0952	0.0587


```

6 -0.0904 -0.0245  0.0512  -0.0164 -0.0338 -0.0658  0.0721  0.203  0.261
7 -0.0161 -0.00185 -0.00131  0.00543  0.0134 -0.0361  0.0145  0.0371 -0.0392
8 -0.712  -0.366  -0.0114  0.0895  -0.163  0.273  -0.435  0.211  -0.0519
9 -0.232  0.463  0.700  0.263  0.380  0.0672 -0.0669  0.0107 -0.0663
10 -0.0402 0.00714 0.0999  -0.0387 -0.0922 -0.509  0.00130 0.303  -0.103
# ... with 29 more rows, and 31 more variables: PC10 <dbl>, PC11 <dbl>,
# PC12 <dbl>, PC13 <dbl>, PC14 <dbl>, PC15 <dbl>, PC16 <dbl>, PC17 <dbl>,
# PC18 <dbl>, PC19 <dbl>, PC20 <dbl>, PC21 <dbl>, PC22 <dbl>, PC23 <dbl>,
# PC24 <dbl>, PC25 <dbl>, PC26 <dbl>, PC27 <dbl>, PC28 <dbl>, PC29 <dbl>,
# PC30 <dbl>, PC31 <dbl>, PC32 <dbl>, PC33 <dbl>, PC34 <dbl>, PC35 <dbl>,
# PC36 <dbl>, PC37 <dbl>, PC38 <dbl>, PC39 <dbl>, feature <chr>

```

```
ggplot(df_out_r, aes(x=PC1, y=PC2, label=feature)) + geom_text_repel()
```

Warning: ggrepel: 32 unlabeled data points (too many overlaps). Consider increasing max.overlaps




4 Lab Exercises

To be handed in via submission of quarto file (and rendered pdf) to GitHub.

1. Using the `delay_2022` data, plot the five stations with the highest mean delays. Facet the graph by line

```
delay_2022 |>
  group_by(line, station_clean) |>
  summarise(mean_delay = mean(min_delay), n_obs = n()) |>
  filter(n_obs>1) |>
  arrange(line, -mean_delay) |>
  slice(1:5) |>
  ggplot(aes(station_clean, mean_delay)) +
  geom_col() +
  coord_flip() +
  facet_wrap(~line, scales = "free_y")
```

``summarise()`` has grouped output by 'line'. You can override using the ``groups`` argument.



2. Using the `opendatatatoronto` package, download the data on mayoral campaign contributions for 2014. Hints:
 - find the ID code you need for the package you need by searching for 'campaign' in the `all_data` tibble above

- you will then need to `list_package_resources` to get ID for the data file
- note: the 2014 file you will get from `get_resource` has a bunch of different campaign contributions, so just keep the data that relates to the Mayor election

```
 ::: {.cell}
```

```
list_package_resources("f6651a40-2f52-46fc-9e04-b760c16edd5c")
```

```
 ::: {.cell-output .cell-output-stdout} # A tibble: 2 x 4      name
format last_mod~1      <chr>                                <chr>                                <chr>
<date>      1 campaign-contributions-2014-data      5b230e92-0a22-4a15-9~
ZIP      2019-07-23  2 campaign-contributions-2014-readme-xls aaf736f4-7468-4bda-9~
XLS      2019-07-23  # ... with abbreviated variable name 1: last_modified
 :::
```

```
all_campaigns <- get_resource("5b230e92-0a22-4a15-9572-0b19cc222985")
```

```
 ::: {.cell-output .cell-output-stderr} ““ New names: New names: New names: New
names: New names: New names: New names:
```

- ‘->...2’
- ‘->...3’

```
 :::
```

```
```.r .cell-code}
df <- all_campaigns[[2]]
```

```
 :::
```

3. Clean up the data format (fixing the parsing issue and standardizing the column names using `janitor`)

```
df <- df |>
 janitor::row_to_names(1) |>
 janitor::clean_names()
```

4. Summarize the variables in the dataset. Are there missing values, and if so, should we be worried about them? Is every variable in the format it should be? If not, create new variable(s) that are in the right format.

```
skim(df)
```

Table 5: Data summary

Name	df
Number of rows	10199
Number of columns	13
Column type frequency: character	13
Group variables	None

**Variable type: character**


skim_variable	n_missing	complete_rate	min	max	empty	n_unique	whitespace
contributors_name	0	1	4	31	0	7545	0
contributors_address	10197	0	24	26	0	2	0
contributors_postal_code	0	1	7	7	0	5284	0
contribution_amount	0	1	1	18	0	209	0
contribution_type_desc	0	1	8	14	0	2	0
goods_or_service_desc	10188	0	11	40	0	9	0
contributor_type_desc	0	1	10	11	0	2	0
relationship_to_candidate	10166	0	6	9	0	2	0
president_business_manager	10197	0	13	16	0	2	0
authorized_representative	10197	0	13	16	0	2	0
candidate	0	1	9	18	0	27	0
office	0	1	5	5	0	1	0
ward	10199	0	NA	NA	0	0	0

```
df <- df |>
 mutate(contribution_amount = as.numeric(contribution_amount))
```

5. Visually explore the distribution of values of the contributions. What contributions are notable outliers? Do they share a similar characteristic(s)? It may be useful to plot the distribution of contributions without these outliers to get a better sense of the majority of the data.

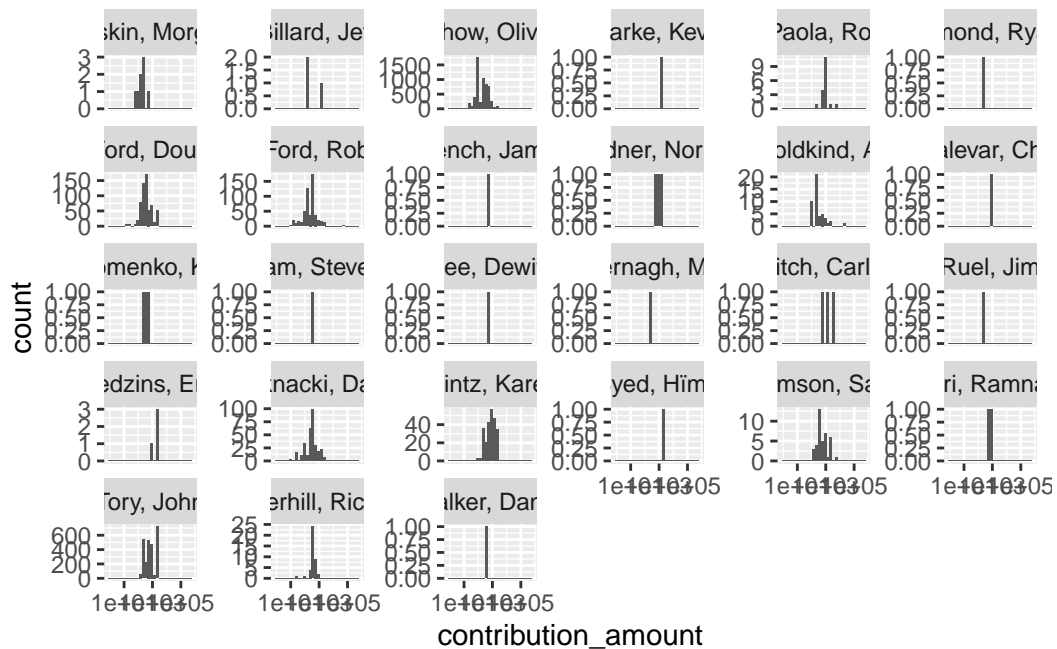
```
df |>
 ggplot(aes(contribution_amount)) + geom_histogram() + scale_x_log10()
```

`stat\_bin()` using `bins = 30`. Pick better value with `binwidth`.



```
df |>
 ggplot(aes(contribution_amount)) + geom_histogram() + scale_x_log10() + facet_wrap(~cand
```

`stat\_bin()` using `bins = 30`. Pick better value with `binwidth`.



```
The big outliers are from Fords to Fords
```

6. List the top five candidates in each of these categories:

- total contributions
- mean contribution
- number of contributions

```
total contributions
df |>
 group_by(candidate) |>
 summarise(total_contr = sum(contribution_amount)) |>
 arrange(-total_contr)
```

```
A tibble: 27 x 2
```

	candidate	total_contr
	<chr>	<dbl>
1	Tory, John	2767869.
2	Chow, Olivia	1638266.
3	Ford, Doug	889897.
4	Ford, Rob	387648.
5	Stintz, Karen	242805



```

6 Soknacki, David 132431
7 Goldkind, Ari 41125.
8 Thomson, Sarah 34628.
9 Di Paola, Rocco 21126
10 Underhill, Richard 15660
... with 17 more rows

```

```

mean contributions
df |>
 group_by(candidate) |>
 summarise(mean_contr = mean(contribution_amount)) |>
 arrange(-mean_contr)

```

```

A tibble: 27 x 2
 candidate mean_contr
 <chr> <dbl>
1 Sniedzins, Erwin 2025
2 Syed, Himy 2018
3 Ritch, Carlisle 1887.
4 Ford, Doug 1456.
5 Clarke, Kevin 1200
6 Di Paola, Rocco 1174.
7 Tory, John 1064.
8 Gardner, Norman 1000
9 Stintz, Karen 995.
10 Kalevar, Chai 900
... with 17 more rows

```

```

number
df |>
 group_by(candidate) |>
 tally() |>
 arrange(-n)

```

```

A tibble: 27 x 2
 candidate n
 <chr> <int>
1 Chow, Olivia 5708
2 Tory, John 2602
3 Ford, Doug 611

```

```

4 Ford, Rob 538
5 Soknacki, David 314
6 Stintz, Karen 244
7 Goldkind, Ari 47
8 Underhill, Richard 41
9 Thomson, Sarah 40
10 Di Paola, Rocco 18
... with 17 more rows

```

7. Repeat 6 but without contributions from the candidates themselves.

```

df_not_to_self <- df |>
 filter(contributors_name!=candidate)

df_not_to_self |>
 group_by(candidate) |>
 summarise(total_contr = sum(contribution_amount)) |>
 arrange(-total_contr)

```

```

A tibble: 17 x 2
 candidate total_contr
 <chr> <dbl>
1 Tory, John 2765369.
2 Chow, Olivia 1634766.
3 Ford, Doug 331173.
4 Stintz, Karen 242805
5 Ford, Rob 174510.
6 Soknacki, David 132431
7 Thomson, Sarah 27702.
8 Goldkind, Ari 17501
9 Underhill, Richard 15660
10 Di Paola, Rocco 15126
11 Ritch, Carlie 5660
12 Sniedzins, Erwin 5600
13 Gardner, Norman 3000
14 Baskin, Morgan 1550
15 Billard, Jeff 1486.
16 Tiwari, Ramnarine 1000
17 Lam, Steven 300

```

```
mean contributions
df_not_to_self |>
 group_by(candidate) |>
 summarise(mean_contr = mean(contribution_amount)) |>
 arrange(-mean_contr)
```

```
A tibble: 17 x 2
 candidate mean_contr
 <chr> <dbl>
1 Ritch, Carlie 1887.
2 Sniedzins, Erwin 1867.
3 Tory, John 1063.
4 Gardner, Norman 1000
5 Tiwari, Ramnarine 1000
6 Stintz, Karen 995.
7 Di Paola, Rocco 890.
8 Thomson, Sarah 729.
9 Ford, Doug 545.
10 Billard, Jeff 496.
11 Soknacki, David 422.
12 Underhill, Richard 382.
13 Goldkind, Ari 380.
14 Ford, Rob 329.
15 Lam, Steven 300
16 Chow, Olivia 286.
17 Baskin, Morgan 194.
```

```
number
df_not_to_self |>
 group_by(candidate) |>
 tally() |>
 arrange(-n)
```

```
A tibble: 17 x 2
 candidate n
 <chr> <int>
1 Chow, Olivia 5706
2 Tory, John 2601
3 Ford, Doug 608
4 Ford, Rob 531
```

5	Soknacki, David	314
6	Stintz, Karen	244
7	Goldkind, Ari	46
8	Underhill, Richard	41
9	Thomson, Sarah	38
10	Di Paola, Rocco	17
11	Baskin, Morgan	8
12	Billard, Jeff	3
13	Gardner, Norman	3
14	Ritch, Carlie	3
15	Sniedzins, Erwin	3
16	Lam, Steven	1
17	Tiwari, Ramnarine	1

8. How many contributors gave money to more than one candidate?

```
df |>
 group_by(contributors_name) |>
 distinct(candidate) |>
 tally() |>
 filter(n>1) |>
 nrow()
```

[1] 184

```
OR

df |>
 group_by(contributors_name, candidate) |>
 tally() |>
 group_by(contributors_name) |>
 tally() |>
 filter(n>1) |> nrow()
```

[1] 184